

YVONNE YU

Phone: 609-373-3202 Email: yy5919@nyu.edu

LinkedIn: [linkedin.com/in/yvonne-yu-76364531b](https://www.linkedin.com/in/yvonne-yu-76364531b)

GitHub: github.com/Yvonne-Yu217

EDUCATION

New York University

M.S. in Data Science, Courant School of Math, Computing & Data Science (GPA: 3.72/4.0) Sept. 2025 – May 2027

Sun Yat-sen University

Bachelor of Economics, Lingnan College (GPA: 3.80/4.0) Sept. 2021 – Jun. 2025

PROFESSIONAL EXPERIENCE

ByteDance

Data Scientist Intern, E-Commerce Search

May 2026 – Aug. 2026

- **Product & Query Research:** Compared placement/query/lifecycle impression–click–purchase funnels with manual demand and q2q/i2q candidates; revealed merchant store/brand-query coverage gaps and appliances/3C cold-start priorities.
- **A/B Experimentation:** Reviewed a historical 12-day test of a 14-day same-category active-revisit ranking term; it lifted mall-entry 30-day active days/user by **0.2026%** ($p < 0.001$) and shelf-channel paid orders/user by **0.1859%** ($p = 0.0063$).
- **Repeat-Search Causal Analysis:** Estimated search/content/mall effects on T+2–T+7 active repeat search via uid-level DML-SparseLinear and 30-day controls; reversal testing found overexposure, guiding recall/exit.
- **Dual-Granularity Query Foundation:** Rebuilt Module/Query datasets in SQL/Spark, unifying placement mappings, impression–click–purchase metrics, quality checks, and trend views for research, experiments, and anomaly drill-down.
- **AI Experiment Automation:** Automated parameter/traffic/Pre-AA/lift/guardrail checks in an AI Skill with multi-agent/human review; supported cross-functional experiment review.

Mars

Data Scientist Intern

May 2025 – Aug. 2025

- **Feature & Candidate Pipeline:** Built incremental resource_id×day ETL for telemetry/metadata/cost; combined 7/30/90-day Pandas/NumPy/Scikit-learn features with Isolation Forest/Z-score/low-utilization rules to flag idle assets.
- **Supervised Ranking & Evaluation:** Trained a second-stage XGBoost ranker on Ops labels with time splits/class weights; at a fixed review budget, **Precision@K** reached about **85% at K=100** in manual review, with asset-level SHAP rationales.
- **Production & Monitoring:** Deployed weekly training/daily batch scoring in Azure ML; Power BI monitored missingness, PSI, score distribution, resource-level precision, and alert volume; backtests/shadow runs gated new SKUs/definitions.
- **Calibration & Governance:** Calibrated anomaly scores to idle probability with Isotonic Regression and set Top-K by Ops capacity; auto-tagged candidates had a 14-day appeal window before ServiceNow ticketing.
- **Feedback & Cost Closure:** Fed confirmed/appealed/closed ServiceNow outcomes into retraining labels; verified decommission/right-sizing yielded annualized savings and cut investigation time from **52 to 31 minutes (-40%)**.

Korn Ferry

Data Science Consultant Intern

Aug. 2024 – Dec. 2024

- **Compensation Data Product:** Built Python/SQL ETL for pay/financial data, unifying currency/year/job family/industry/region in a traceable query layer; cut annual pay review from **two weeks to two days (-80%)**.
- **Talent Segmentation:** Standardized 28 pay/performance/tenure features across employee records; PCA reduced them to 6 dimensions and K-Means formed **4 clusters**, yielding profiles for retention, succession, development, and job pricing.

RESEARCH & PROJECTS

Retrieval-Coordinated Reasoning for Small LLMs

Feb. 2026 – May 2026

- Evaluated **7 models and 10 strategies across 4 datasets** (uneven coverage); paired tests found voting weaker than the best single model; diagnostics suggested correlated errors and miscalibration.
- Built RnA1/F1 RAG pipelines; F1 lifted ConTRoL accuracy **65.5% to 71.7% (+6.2 pp)** vs. Qwen3-8B single-model retrieval; LettuceDetect flagged **< 0.02%** of retrieved-evidence characters as unsupported (805 cases).

Seeing Sentiment: News Images, Retail Participation & Chinese A-Share Returns

Nov. 2024 – Jun. 2026

- Built NLP/CV multimodal sentiment: fine-tuned ViT-B/16 via PyTorch/HuggingFace transfer learning (**88.1%** PCNN accuracy); scored 572K images and encoded 245K articles with Chinese RoBERTa; MongoDB managed a 12-year corpus.
- Ran Newey–West regressions in R, placebo tests, and expanding-window OOS; replicated in statsmodels with past-only scaling; found incremental high-volatility signal and ChiNext $R^2_{OOS} = \mathbf{0.293\%}$ (Clark–West $p = 0.015$).